

RDFQuotient documentation

March 19, 2019

1 RDFQuotient

RDFQuotient tool, hosted at Inria GitLab repository¹, is an implementation of the framework for computing quotient summaries of RDF graphs². Current version of the software: 1.8-SNAPSHOT. A link to download a stand-alone jar file and building-from-sources instructions can be found in README file³. RDFQuotient uses PostgreSQL DBMS as a back-end and it is a requirement for the setup. PostgreSQL server must be running with the specified configuration.

2 Command-line, stand-alone application interface

The command-line interface consists of three possible operations: `load`, `summarize` and `read`. Even though `read` operation is present in command-line interface, it doesn't have any effect and was added for completeness. It is meant to be used in a programmatic way. The operations follow a specific syntax that can be seen in the RDFQuotient manual in Figure 1.

Specifying `dataset.filename` argument for `--load` operation is the only way to change the input graph filename. The `--summarize` operation has two possible settings: `dataset.filename` and `database.name`. Specifying one of them is mandatory. However if both of them are present `database.name` takes precedence. The reason for keeping the two options is that sometimes it is more convenient to use the same name for the argument, especially when we rely on the default naming of the databases created for summaries.

Example command-line execution for loading:

```
java -jar target/RDFQuotient-1.8-SNAPSHOT-with-dependencies.jar --load  
"dataset.filename=/data/datasets/enelshops/enelshops.nt"
```

Example command-line execution for summarization in a dry-run mode:

```
java -jar target/RDFQuotient-1.8-SNAPSHOT-with-dependencies.jar --summarize  
"dataset.filename=/data/datasets/enelshops/enelshops.nt" --dry-run
```

Example command-line execution for summarization:

```
java -jar target/RDFQuotient-1.8-SNAPSHOT-with-dependencies.jar --summarize  
"dataset.filename=/data/datasets/enelshops/enelshops.nt"
```

2.1 Configuration

The tool encapsulates hard-coded values for configuration that model all the available functionalities and scenarios. They can be overwritten by command-line arguments or by two configuration files: `loading.properties` and `summarization.properties`. The command-line arguments and configuration files are optional. The configuration files serve a purpose of templates, they are filled in with the default, hard-coded values.

¹<https://gitlab.inria.fr/cedar/RDFQuotient>

²<https://project.inria.fr/rdfquotient/>

³<https://gitlab.inria.fr/cedar/RDFQuotient/blob/master/README.md>

usage: RDFQuotient 1.8-SNAPSHOT

This framework is designed to work with one graph at a time.
Before using RDFQuotient make sure that Postgres server is running.
Input RDF dataset file format is N-Triples and the file is assumed not to
contain any duplicated triples.

```
rdfquotient [-d] [-h | -l <[ARGS]> | -r <[ARGS]> | -s <[ARGS]> | -v] [-lp  
               <filename> | -sp <filename>]  
-d,--dry-run                - print the configuration used for a  
                             run  
-h,--help                  - print help message  
-l,--load <[ARGS]>         - load the dataset into database  
                             [ARGS] must specify a value for  
                             dataset.filename  
                             CAUTION: database is dropped by  
                             default  
-lp,--loading-properties <filename> - set loading properties filename  
-r,--read <[ARGS]>         - read summary from database  
                             [ARGS] must specify a value for  
                             database.name  
-s,--summarize <[ARGS]>    - summarize the dataset  
                             [ARGS] must specify a value for  
                             dataset.filename or database.name  
-sp,--summarization-properties <filename> - set summarization properties  
                             filename  
-v,--version               - print the version of RDFQuotient
```

[ARGS] is a comma-separated list of assignments of form key=value, where key is a
configuration property from the list of loading or summarization configuration
properties.

Figure 1: Command-line interface of RDFQuotient

To determine the configuration for a run of the tool for either loading or summarization properties files, Algorithm 1 applies. Hard-coded values have the lowest priority. Values present in the configuration files have medium priority. Values passed in the arguments have the highest priority.

Algorithm 1 Setting up of the configuration

```
1: procedure SETUPCONFIGURATION  
2:   if args["properties_filename"]  $\neq$  null then  
3:     properties_filename  $\leftarrow$  args["properties_filename"]  
4:   if properties_filename file exists then  
5:     overwrite default, hard-coded values for properties with the values in properties_filename file  
6:   for each setting  $\in$  args.values do  
7:     overwrite value of setting in properties with the value specified by args[setting]
```

2.2 Configuration files

The following two lists contain all the functionalities and scenarios in terms of the configuration parameters of the execution.

2.2.1 Loading properties

The list of loading configuration properties with default values assigned:

1. `dataset.filename = test.nt`
Caution: RDFQuotient tool assumes the N-Triples format for the dataset file and no duplicates present.
2. `database.host = localhost`
3. `database.port = 5432`
4. `database.user = postgres`
5. `database.password = postgres`
6. `database.name =`
Blank by default.
If left blank, database name will be derived from dataset filename. Otherwise, the specified name will be used.
7. `database.drop_existing_db = true`
Valid options: `true`, `false`.
Caution: if `database.name` exists, it is dropped by default before loading.
8. `database.storage_layout = TRIPLES_TABLE`
Valid options: `TRIPLES_TABLE`, `TABLE_PER_ROLE_AND_CONCEPT`.
9. `database.triples_table_name = triples`
10. `database.encoded_triples_table_name = encoded_triples`
11. `database.encoded_saturated_triples_table_name = encoded_saturated_triples`
12. `database.dictionary_table_name = dictionary`
13. `dictionary.fetch_size = 1000`
14. `saturation.enable = false`
Valid options: `true`, `false`.
15. `database.encoded_saturated_triples_table_name = encoded_saturated_triples`
16. `saturation.batch_size = 1000`
17. `database.deterministic_ordering = false`
Whether to sort dictionary entries before assigning encodings and `alsoencode_triples` and `encoded_saturated_triples` tables. Used for tests, should not be used for big graphs.
18. `statistics.export_to_csv_file = true`
Valid options: `true`, `false`.
19. `configuration.export_to_disk = false`
Whether to export `loading.properties` used in a run after it's finished.
Valid options: `true`, `false`.

2.2.2 Summarization properties

The list of summarization configuration properties with default values assigned:

1. `dataset.filename = test.nt`
2. `database.host = localhost`
3. `database.port = 5432`
4. `database.user = postgres`
5. `database.password = postgres`
6. `database.name =`
Blank by default.
If left blank, database name will be derived from dataset filename. Otherwise, the specified name will be used.
7. `database.triples_table_name = triples`
8. `database.encoded_triples_table_name = encoded_triples`
9. `database.dictionary_table_name = dictionary`
10. `database.encoded_saturated_triples_table_name = encoded_saturated_triples`
11. `summary.summarize_saturated_graph = false`
Raises an error in case of absence of saturated triples table.
Valid options: `true`, `false`.
12. `summary.type = typedweak`
Valid options: `weak`, `strong`, `typedweak`, `typedstrong`, `2pweak`, `2pweakunionfind`, `2pstrong`, `2ptypedweak`, `2ptypedstrong`, `onefb`, `onefw`.
13. `summary.omit_generic_properties_from_cliques = true`
Valid options: `true`, `false`.
14. `summary.generic_properties = <http://www.w3.org/2000/01/rdf-schema#label>,
<http://www.w3.org/2000/01/rdf-schema#comment>,
<http://www.w3.org/1999/02/22-rdf-syntax-ns#value>,
<http://www.w3.org/2002/07/owl#sameAs>`
List of comma-separated values of URI within triangular brackets.
15. `summary.replace_type_with_most_general_type = true`
Raises an error in case of W and S summaries.
Valid options: `true`, `false`.
16. `summary.consistency_checks = false`
This may make summarization significantly slower; set it to `true` only for debugging.
Valid options: `true`, `false`.
17. `summary.export_to_nt_file = true`
Valid options: `true`, `false`.
18. `summary.nt_file_prefix = summariesNT/`
Prefix added to summary NT file, usually to dispatch it into a separate directory.
19. `summary.export_to_database = true`
Valid options: `true`, `false`.

20. `summary.save_representation_function_and_node_statistics = true`
Valid options: `true`, `false`.
21. `summary.gather_representation_counts = true`
Valid options: `true`, `false`.
22. `summary.step_by_step = false`
Whether to execute step-by-step summarization. It is recommended to set `drawing.step_by_step = true`.
23. `statistics.export_to_csv_file = true`
Valid options: `true`, `false`.
24. `configuration.export_to_disk = false` Whether to export summarization.properties used in a run after it's finished.
Valid options: `true`, `false`.
25. `drawing.dot_installation = /usr/local/bin/dot`
Path to dot executable. If this is not found, it is not an error, just drawing won't work.
26. `drawing.remove_dot_file = false`
Whether to remove DOT file in case of successful execution of dot.
27. `drawing.dot_file_prefix = summariesDOT/`
Prefix added to summary DOT file, usually to dispatch it into a separate directory.
28. `drawing.png_file_prefix = summariesPNG/`
Prefix added to summary PNG file, usually to dispatch it into a separate directory.
29. `drawing.step_by_step = false`
Whether to draw step-by-step drawings.
30. `drawing.style = split_and_fold_leaves`
Valid options: `plain`, `split_leaves`, `split_and_fold_leaves`.
If an invalid options is specified, drawing will be suppressed.
31. `drawing.color_scheme = diverse`
Valid options: `diverse`, `bw`, `ivory`, `light`.
32. `drawing.summary_node_URI_prefix = http://rq.org/`
33. `drawing.title = true`
Valid options: `true`, `false`.
34. `drawing.max_node_label_length = 20`
35. `drawing.max_types_drawn_per_namespace = 5`
36. `drawing.summary_node_support_URI_prefix = http://rq.org/nodeSupport`
37. `drawing.summary_edge_support_URI_prefix = http://rq.org/edgeSupport`
38. `drawing.reified_summary_edge_URI_prefix = http://rq.org/edge`
39. `drawing.reified_edge_subject_URI_prefix = http://rq.org/reifiedEdgeSubject`
40. `drawing.reified_edge_property_URI_prefix = http://rq.org/reifiedEdgeProperty`
41. `drawing.reified_edge_object_URI_prefix = http://rq.org/reifiedEdgeObject`

3 Programmatic Java interface

The same set of options, as for stand-alone interface, is also available for programmatic Java interface: `fr.inria.cedar.RDFQuotient.controller.Interface` class.

Here, instead of passing command-line arguments to the program, we pass a `Properties` Java object to one of dedicated methods. We can explicitly set values of some configuration properties using other dedicated methods. We can also access the `Connection` Java object and not close it if we don't want to. It makes this implementation stateful as it can potentially use a connection opened outside of this class and can return it to be used later. The command-line implementation is stateless.

See detailed explanation of methods in `Interface` class below:

1. `public static Connection getOrEstablishNewDatabaseConnection(Properties properties)`
If the database connection is established it returns it, otherwise establishes new connection using configuration specified in `properties`.
2. `public static Connection getDatabaseConnection()`
Returns database connection, may return `null` if database connection was not established or closed.
3. `public static void closeDatabaseConnection()`
Closes database connection.
4. `public static void forceCloseDatabaseConnection(Properties properties)`
This is an emergency method if some library misbehaves and doesn't close connection, should not be used otherwise.
5. `public static String load(String configurationFilename, Properties properties, boolean closeConnection)`
If both `configurationFilename` and `properties` are null or point to a file/object that does not contain valid configuration properties, default values will be used.
Otherwise the following procedure applies:
 - (a) Take default values
 - (b) Overwrite them with the values from the file `configurationFilename`
 - (c) Overwrite them with the values from the object `properties`

Parameter `closeConnection` controls whether to close connection to the database after the call to this method.

Examples:

- (a)

```
Properties myProperties = new Properties();
myProperties.put("dataset.filename", "datasets/my_dataset.nt");
Interface.load>LoadingProperties.getDefaultLoadingPropertiesFilename(),
myProperties, true);
```
- (b)

```
Properties myProperties = LoadingProperties.getDefaultProperties();
myProperties.put("dataset.filename", "datasets/my_dataset.nt");
Interface.load(null, myProperties, true);
```
- (c)

```
Properties myProperties = LoadingProperties.getDefaultProperties();
myProperties.put("dataset.filename", "datasets/my_dataset.nt");
Interface.load("conf/my_config_file.properties", myProperties, true);
```

Returns a name of the database.

6. `public static HashMap<String, String> summarize(String configurationFilename, Properties properties, boolean closeConnection)`
See load method comment: in examples use `SummarizationProperties` class instead of `LoadingProperties`.
Returns a map with keys: `"databaseName"`, `"NTFilename"` and `"DOTFilename"`. If the corresponding name is not present, the value is set to `null`.

7. `public static Summary read(String configurationFilename, Properties properties, boolean closeConnection)`
See load method comment.
Returns Summary object.
8. `public static HashMap<String, String> loadAndSummarize(Properties loadingProperties, Properties summarizationProperties)`
For convenience, this method combines load and summarize functionalities in one. By default the database connection is closed.
9. `public static HashMap<String, String> loadAndSummarizeThroughShortcut(Properties loadingProperties, Properties summarizationProperties)`
For convenience, this method combines load and summarize through shortcut functionalities in one. By default the database connection is closed.

4 Tests

Tests are stored in `src/test` directory. Resources are stored in `src/test/resources` directory that contains all the input test files for our automatic correctness tests (not to confuse with `consistencyChecks` which are tests that can be performed during the execution of the program depending on the configuration). They are checked once `mvn install` command is run. They can be skipped by adding a flag `-DskipTests`, e.g. running a command `mvn clean install -DskipTests`.

5 Scripts

Project tree contains `scripts` directory with useful bash scripts for computing the summaries with different configurations specified. These scripts also account for memory settings of Java virtual machine, as well as, they can realize shortcut computations using stand-alone interface by composing a pipeline of commands.